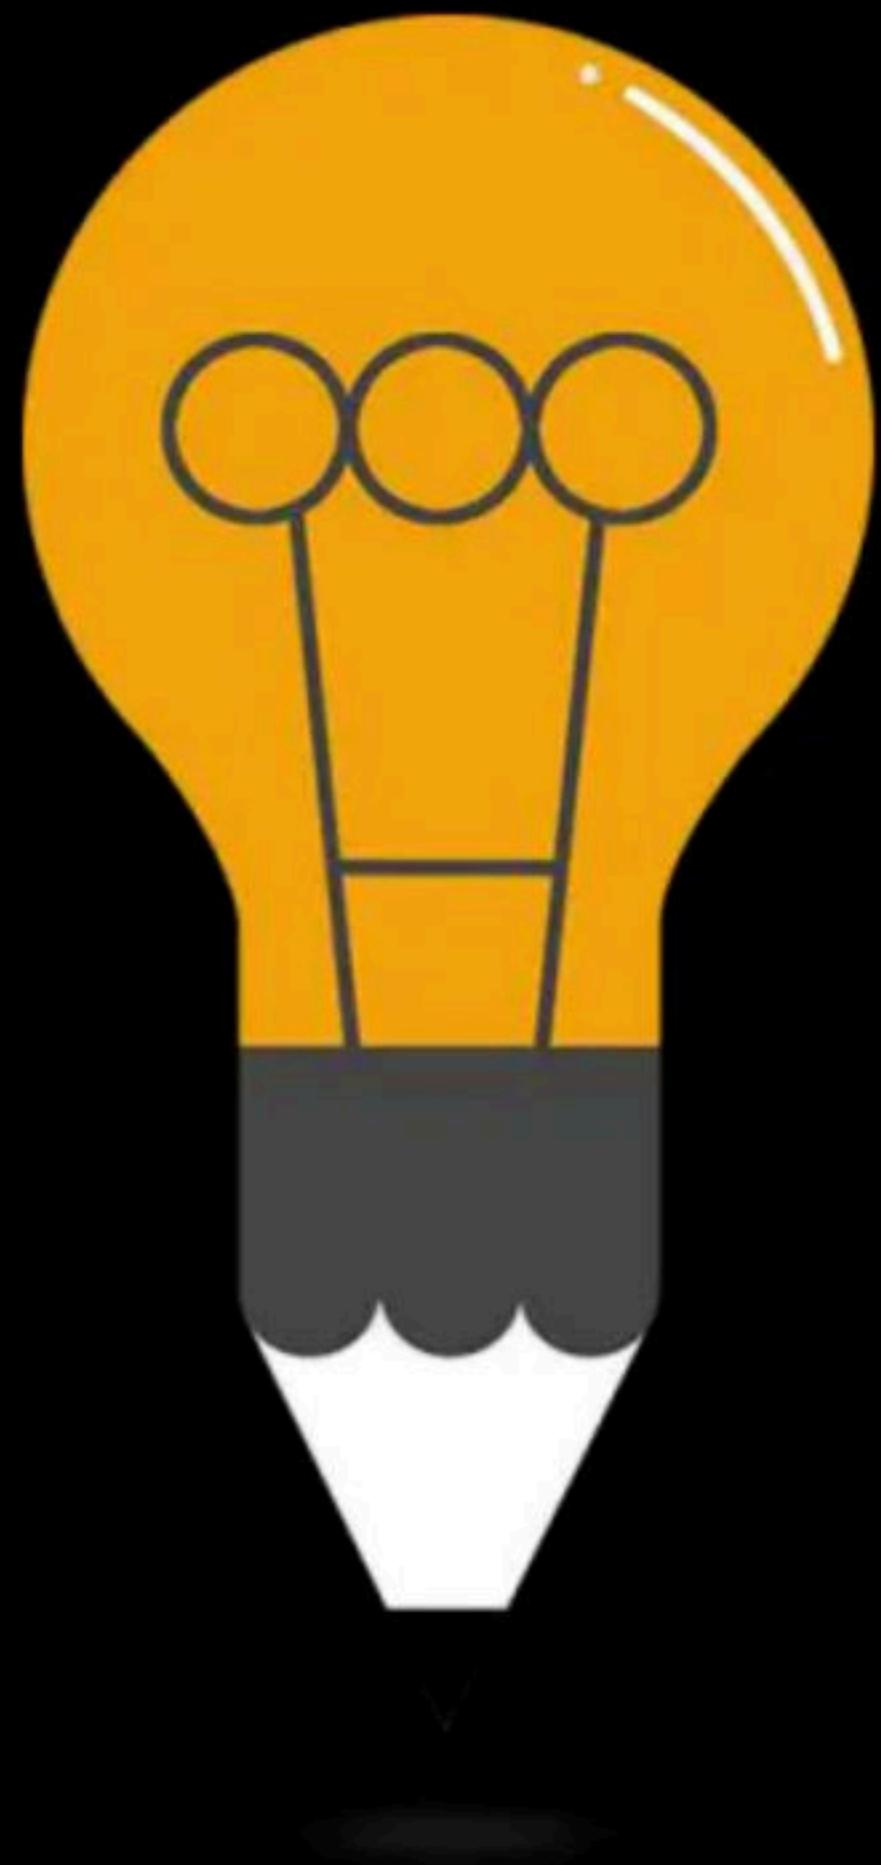


File Organization and Indexing: Part I

Complete Course on Database Management System

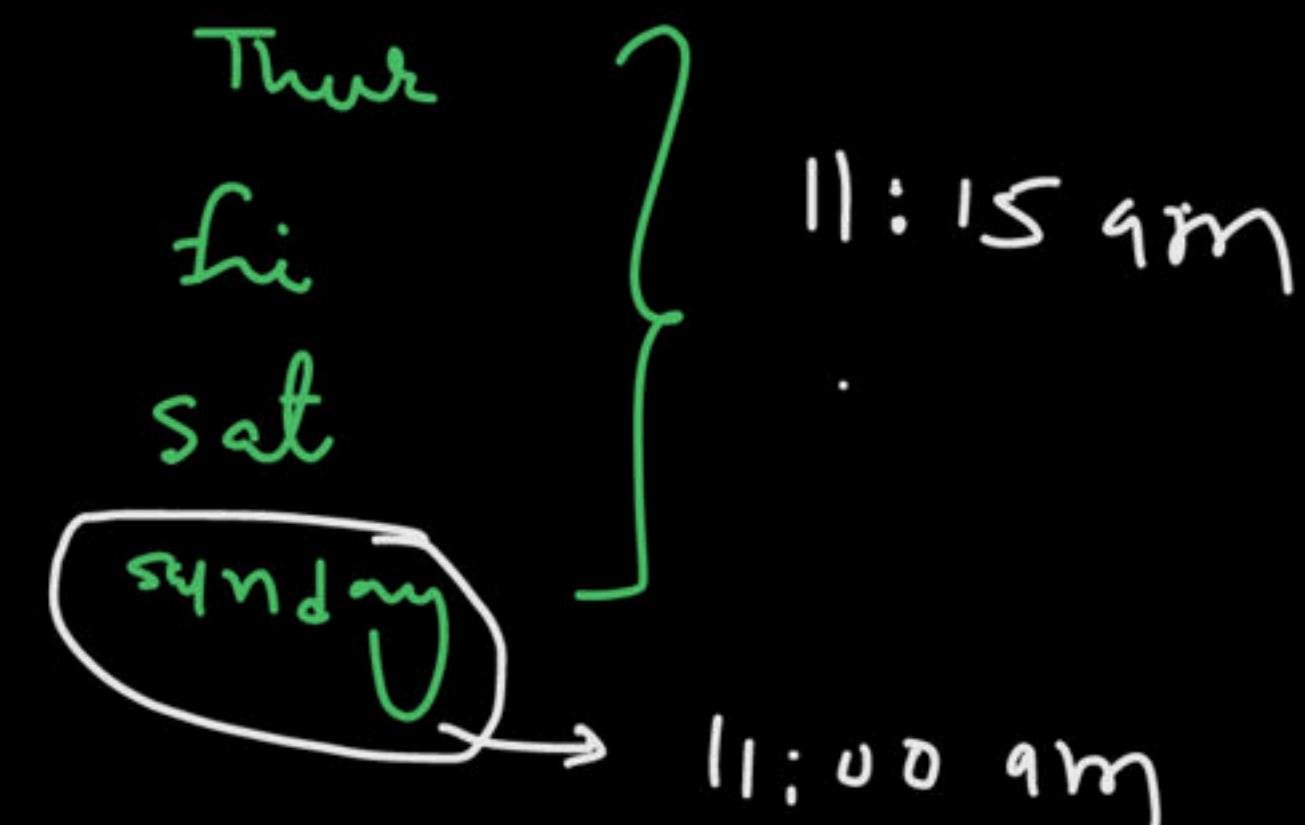


DBMS

Timestamp based protocols & Indexing

By: Vishvadeep Gothi

16, 17



Today ⇒ 2 sessions
11:15 am } only today
11:00 pm }

▲ 1 • Asked by Vaishnavij...

sir in this example if we are not considering the effect of $W(x)$ by T_1 , then the final value should be 25, how 30?

Basic Timestamp Algorithm

$R - TS(x) = \cancel{10}$

$W - TS(x) = \cancel{15} \cancel{20}$

$x = \cancel{15} \cancel{20} 30$

$\frac{T_1}{R(x)}$
 $w(x)$

$\frac{T_2}{w(x)}$
 $w(x)$

$\frac{T_3}{w(x)}$

▲ 1 • Asked by Kumar

Please help me with this doubt

gate syllabus 2023.pdf 3. ER_Modeling_Part_II_... +

file:///D:/Study%20Material/dbms/3.%20ER_Modeling_Part_II_with_Anno.pdf

28 of 47 ⌂ Fit to page Page view Read aloud Add notes

and E2 participate totally in R and that the cardinality of E1 is greater than the cardinality of E2.

Which one of the following is true about R?

A Every entity in E1 is associated with exactly one entity in E2
B Some entity in E1 is associated with more than one entity in E2
C Every entity in E2 is associated with exactly one entity in E1
D Every entity in E2 is associated with at most one entity in E1

↳ no. of entities in entity set



Can I have nwm at E1 for E2

unacademy

A yellow hand-drawn arrow points from the question "Can I have nwm at E1 for E2" towards the relationship line in the ER diagram.

▲ 2 • Asked by Shreyas

Please help me with this doubt

Consider the following log sequence of two transactions on a bank account, with initial balance 12000, that transfer 2000 to a mortgage payment and then apply a 5% interest.

1. T1 start
2. T1 B old = 12000 new = 10000
3. T1 M old = 0 new = 2000
4. T1 commit
5. T2 start
6. T2 B old = 10000 new = 10500
7. T2 commit

Suppose the database system crashes just before log record 7 is written. When the system is restarted, which one statement is true of the recovery procedure?

- A. We must redo log record 6 to set B to 10500
- B. We must undo log record 6 to set B to 10000 and then redo log records 2 and 3
- C. We need not redo log records 2 and 3 because transaction T1 has committed
- D. We can apply redo and undo operations in arbitrary order because they are idempotent

▲ 1 • Asked by Divya

Sir wait_&_die mein younger one(if requests) will get aborted and will restart with "same timestamp"? sir yeh same timestamp mtlb? same arrival time to nhi hoga as it was before?

▲ 1 • Asked by Faizan

Please help me with this doubt

[2014 (Set-3) : 2 Marks]

Q1.55 Consider a uniprocessor system executing three tasks T_1 , T_2 and T_3 , each of which is composed of an infinite sequence of jobs (or instances) which arrive periodically at intervals of 3, 7 and 20 milliseconds, respectively. The priority of each task is the inverse of its period, and the available tasks are scheduled in order of priority, with the highest priority task scheduled first. Each instance of T_1 , T_2 and T_3 requires an execution time of 1, 2 and 4 milliseconds, respectively. Given that all tasks initially arrive at the beginning of the 1st millisecond and task preemptions are allowed, the first instance of T_3 completes its execution at the end of _____ milliseconds.

[2015 (Set-1) : 2 Marks]

159

GS PyQ vishvadeep

Timestamp

Read Timestamp(A): Youngest transaction who read A

Write Timestamp(A): Youngest transaction who write A

Basic Timestamp Algorithm

(Timestamp ordering)

Whenever a Transaction T issues a $W_{item}(X)$ operation, check the following conditions:

- If $R_{TS}(X) > TS(T)$ or if $W_{TS}(X) > TS(T)$, then abort and rollback T and reject the operation.
- Else execute $W_{item}(X)$ operation of T and set $W_{TS}(X)$ to $TS(T)$.

Whenever a Transaction T issues a $R_{item}(X)$ operation, check the following conditions:

- If $W_{TS}(X) > TS(T)$, then abort and rollback T and reject the operation, else
 - If $W_{TS}(X) \leq TS(T)$, then execute the $R_{item}(X)$ operation of T and set $R_{TS}(X)$ to the larger of $TS(T)$ and current $R_{TS}(X)$.
-

Basic Timestamp Algorithm

Restarted transaction gets a younger timestamp

→ eliminates starvation

Basic Timestamp Algorithm

T1 T2
R(X) or W(X)

T1 T2
W(X)

W(X) R(X)

Basic Timestamp Algorithm

T1 T2 T3
R(X)
R(X)
R(X)
W(X)

Question

T1 T2 T3 T4

R(X)

R(X)

R(X)

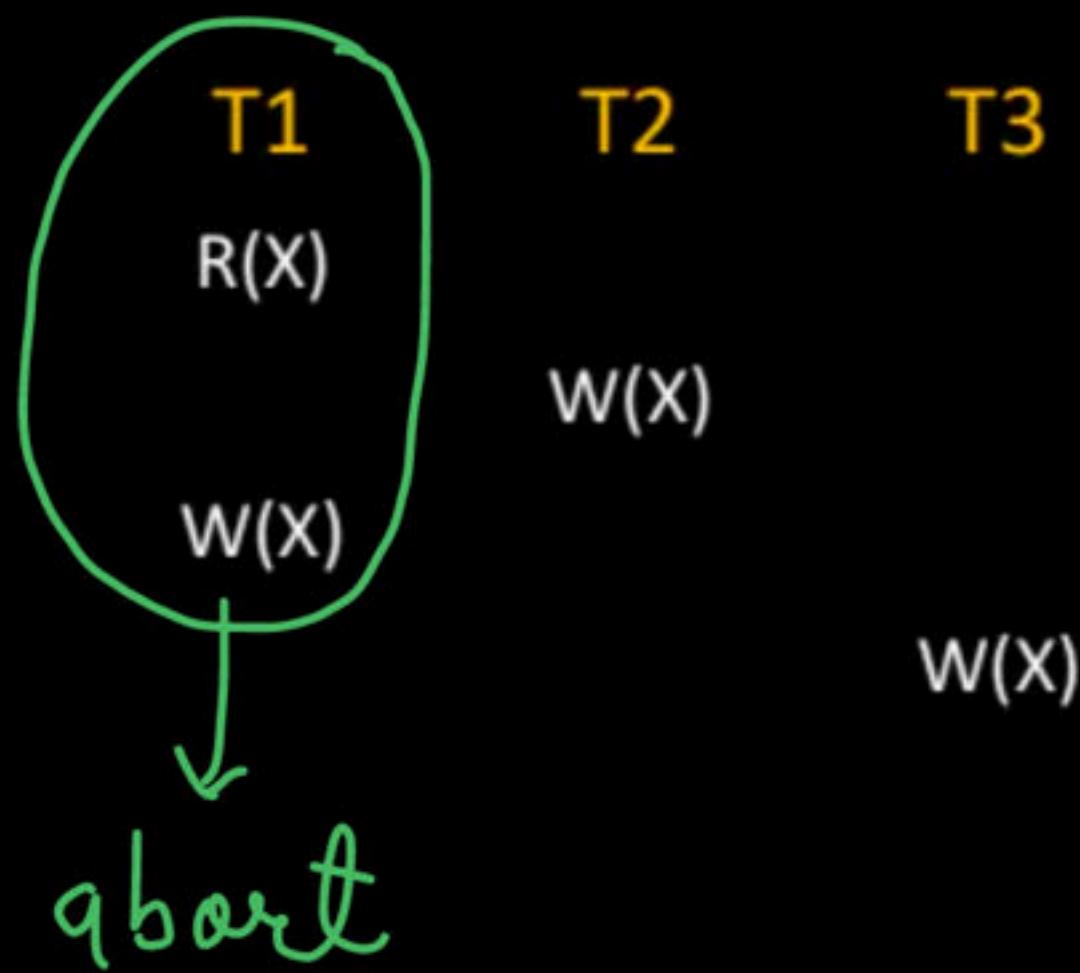
W(X)

W(X)

W(X)

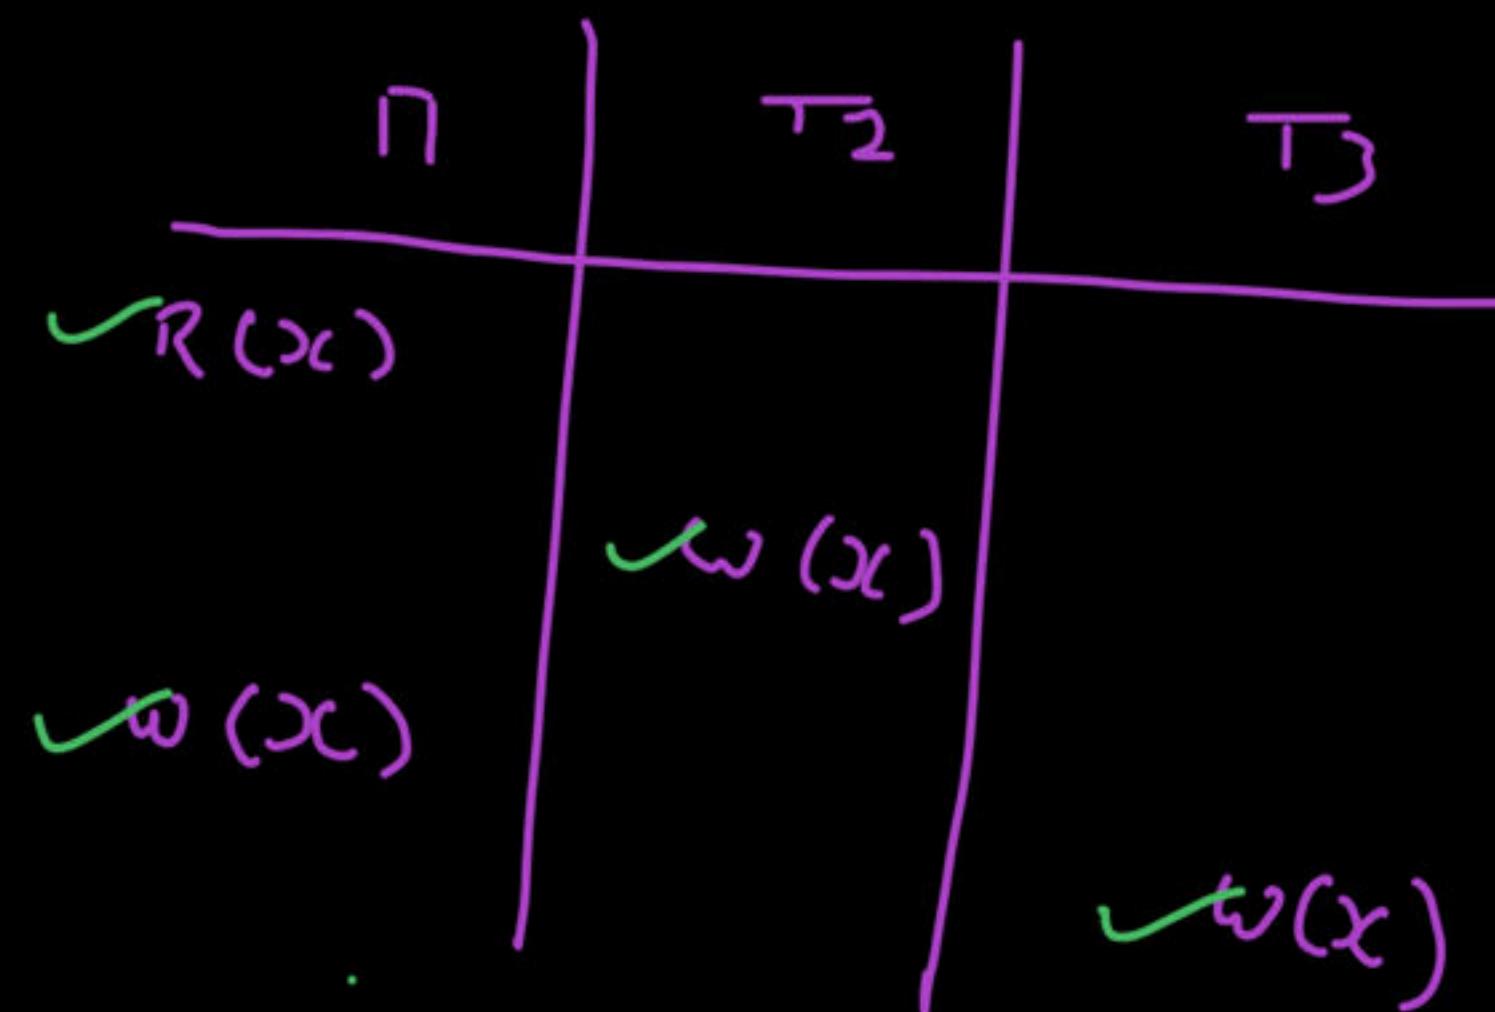
R(X)

Basic Timestamp Algorithm



Thomas Write Rule

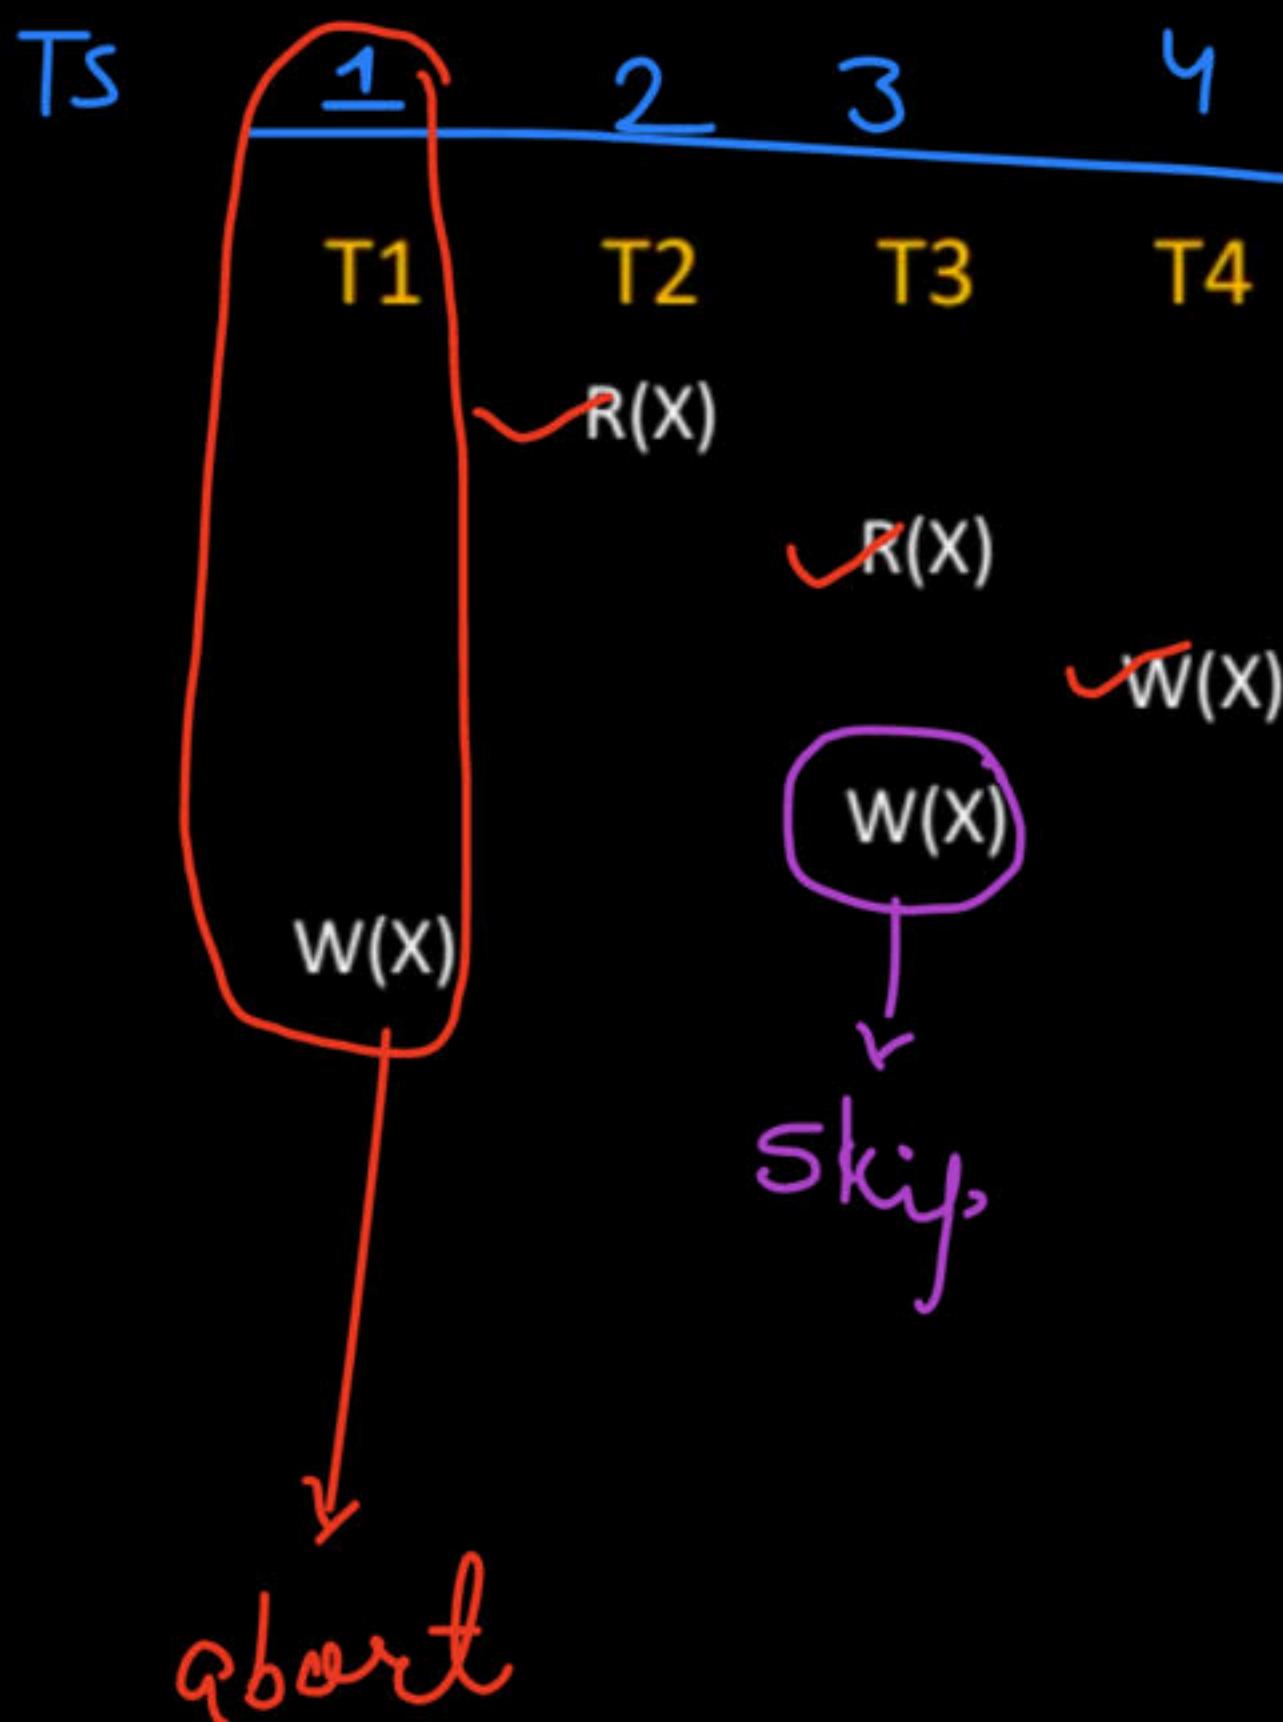
- Read is same as basic timestamp rules
- Write(A) in transaction T:
 - If $R_{TS}(A) > TS(T)$ then abort T, rollback and restart T
 - Else If $W_{TS}(A) >$ then skip write operation $W_{TS}(A) > TS(T)$
 - Else perform Write(A) of T and update $W_{TS}(A) = TS(T)$



$$R_{TS}(x) = 0, 1$$

$$w_{TS}(x) = 0, 2, 3$$

Question



$$R - TS(x) = \sigma \not\leq 3$$

$$W - TS(x) = \emptyset \neq 4$$

$\tau_1 \quad \tau_2 \quad \tau_3 \quad \tau_4$

$\checkmark R(x)$

$\omega(x)$

$\checkmark R(x)$

$\omega(x)$

$R_{\neg TS}(x) \neq \emptyset$ x_2

$\omega_{\neg TS}(x) \neq \emptyset$ x_3, x_4

$\omega(x)$

skip

$\omega(x)$

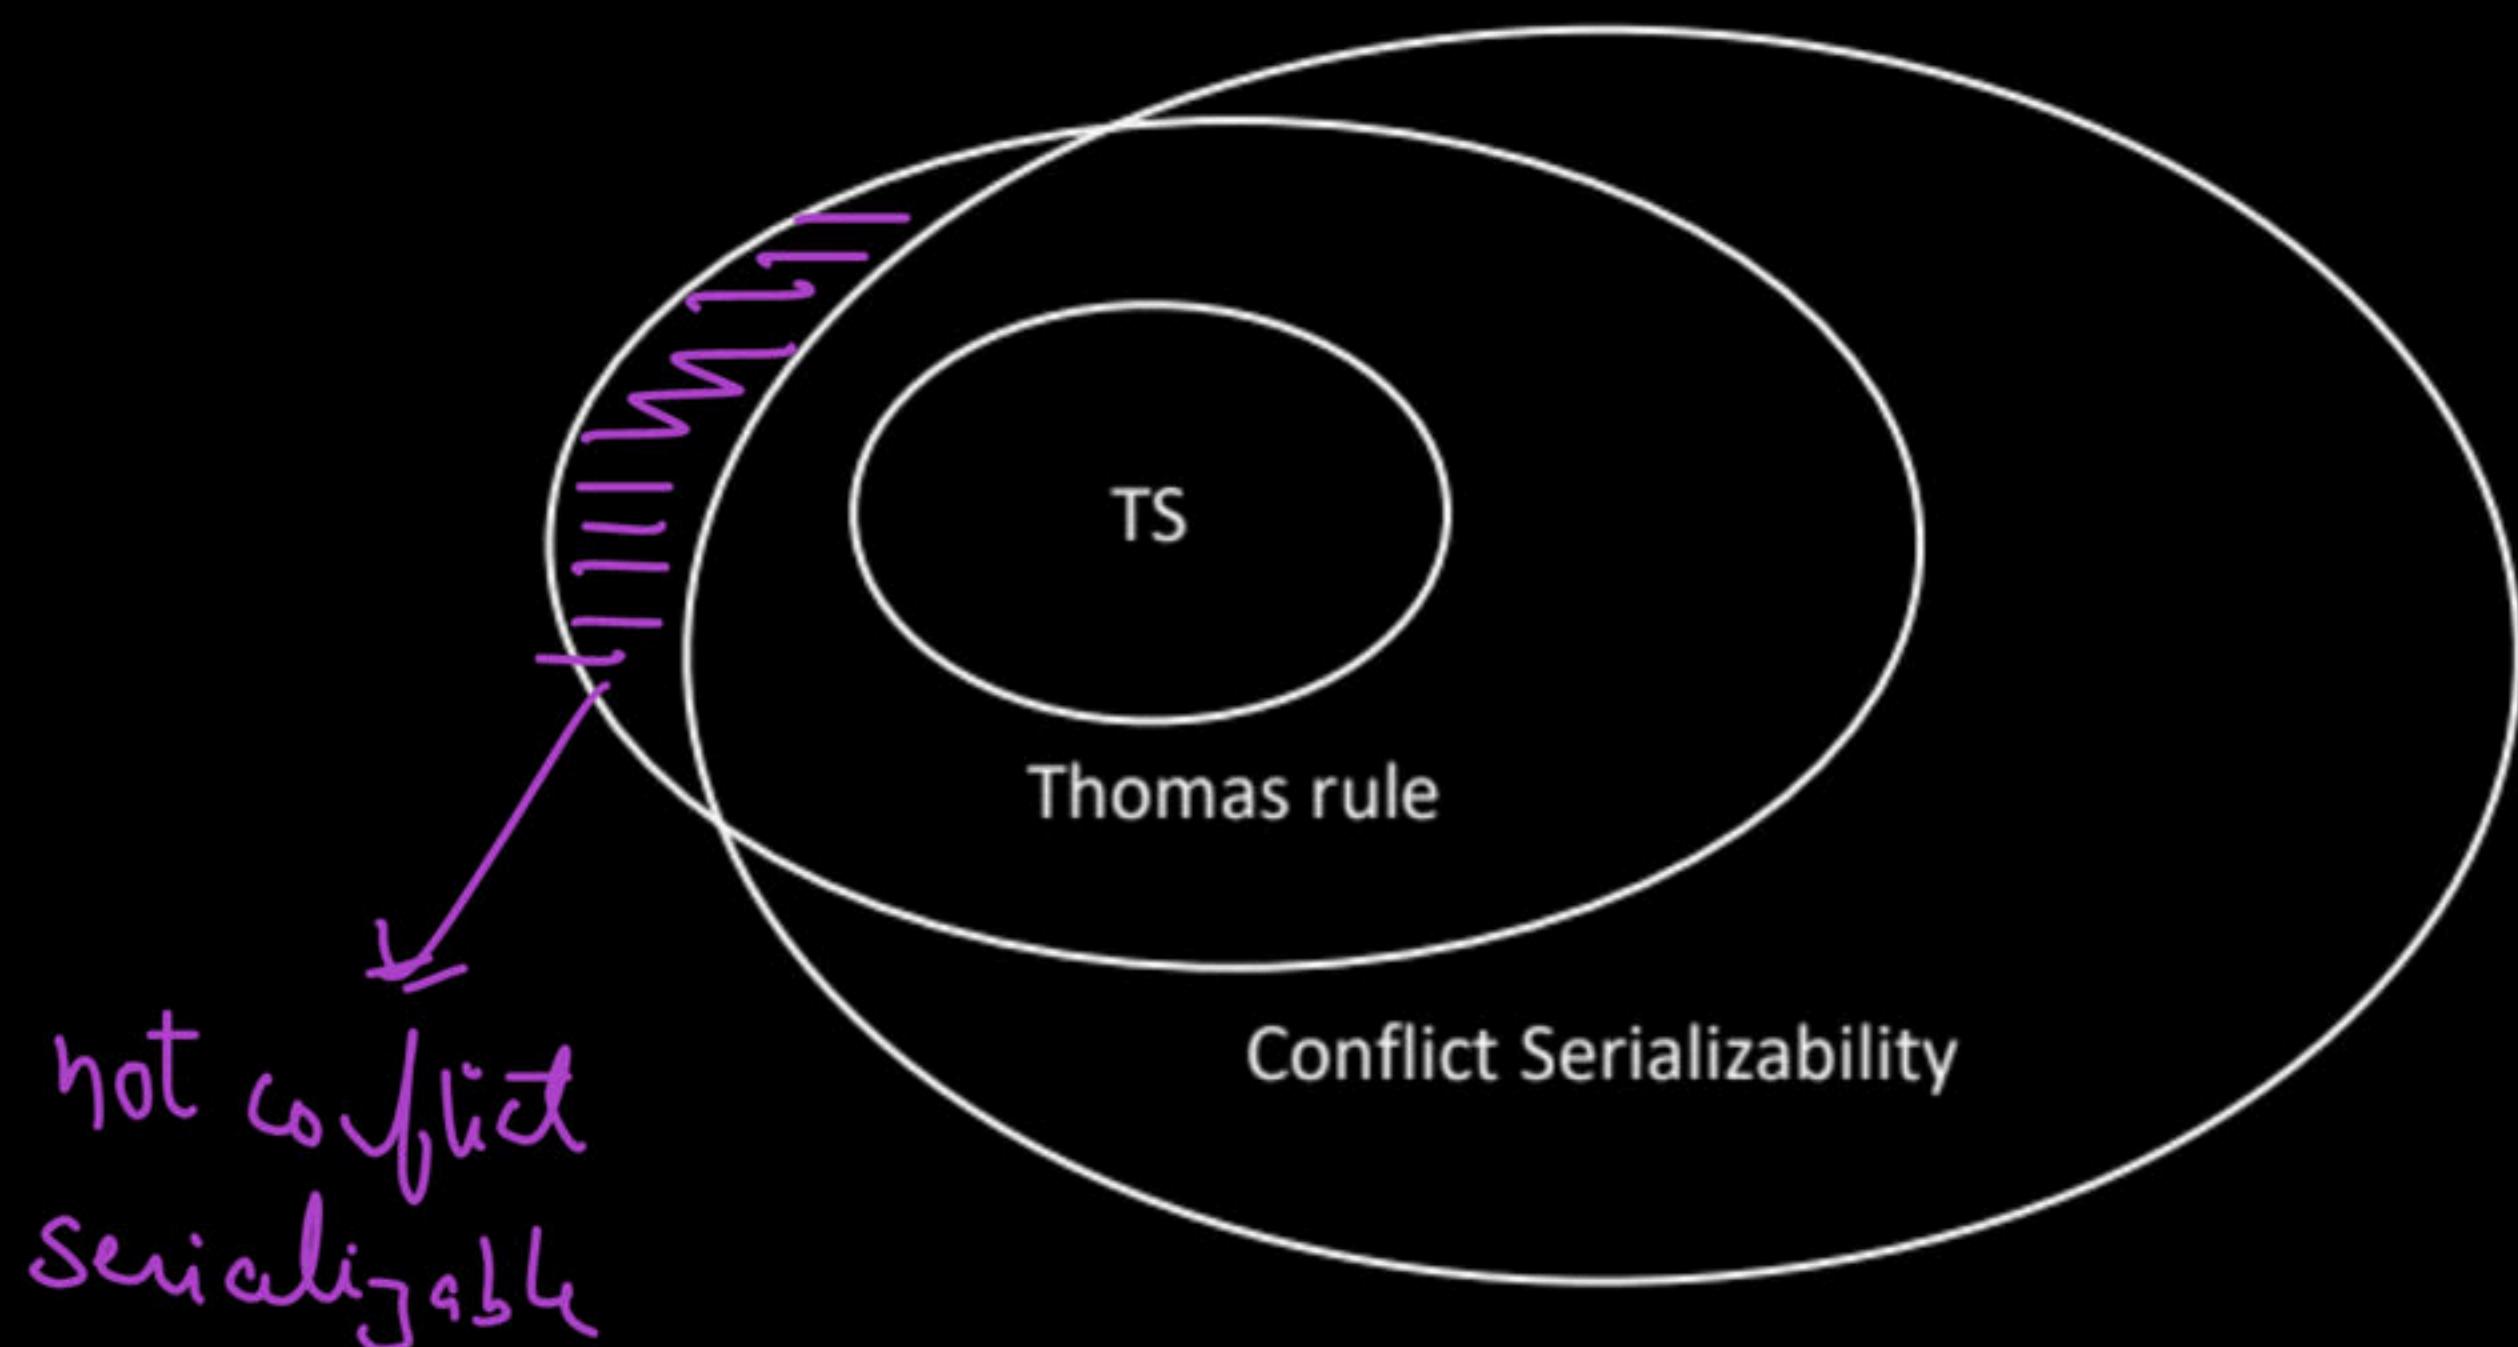
$\omega(x)$ → skip

Timestamp ordering alg^b

- serializability
- No deadlock

Timestamp

- Basic TS allows conflict serializable schedules
- Thomas rule allows more than conflict serializable schedules

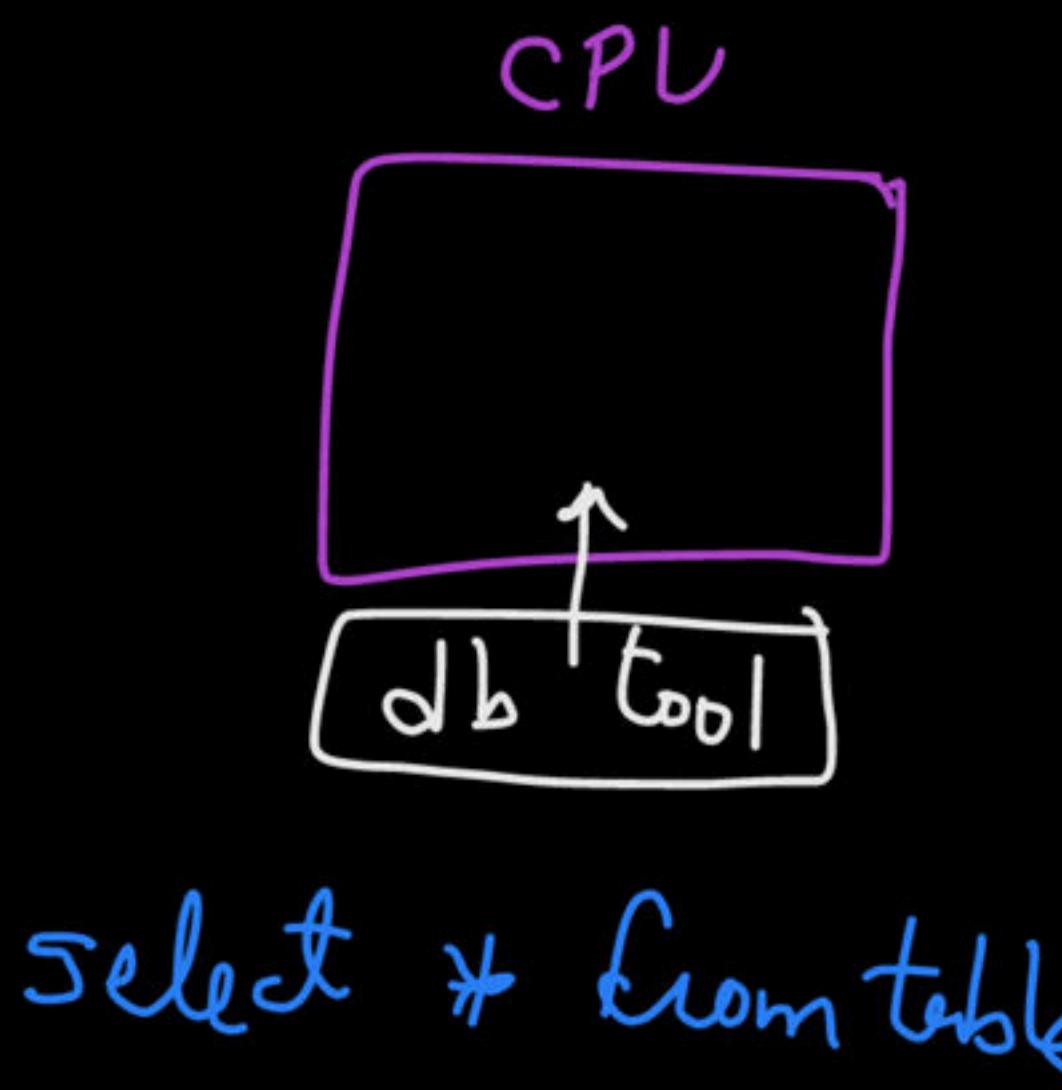


No Need to Study

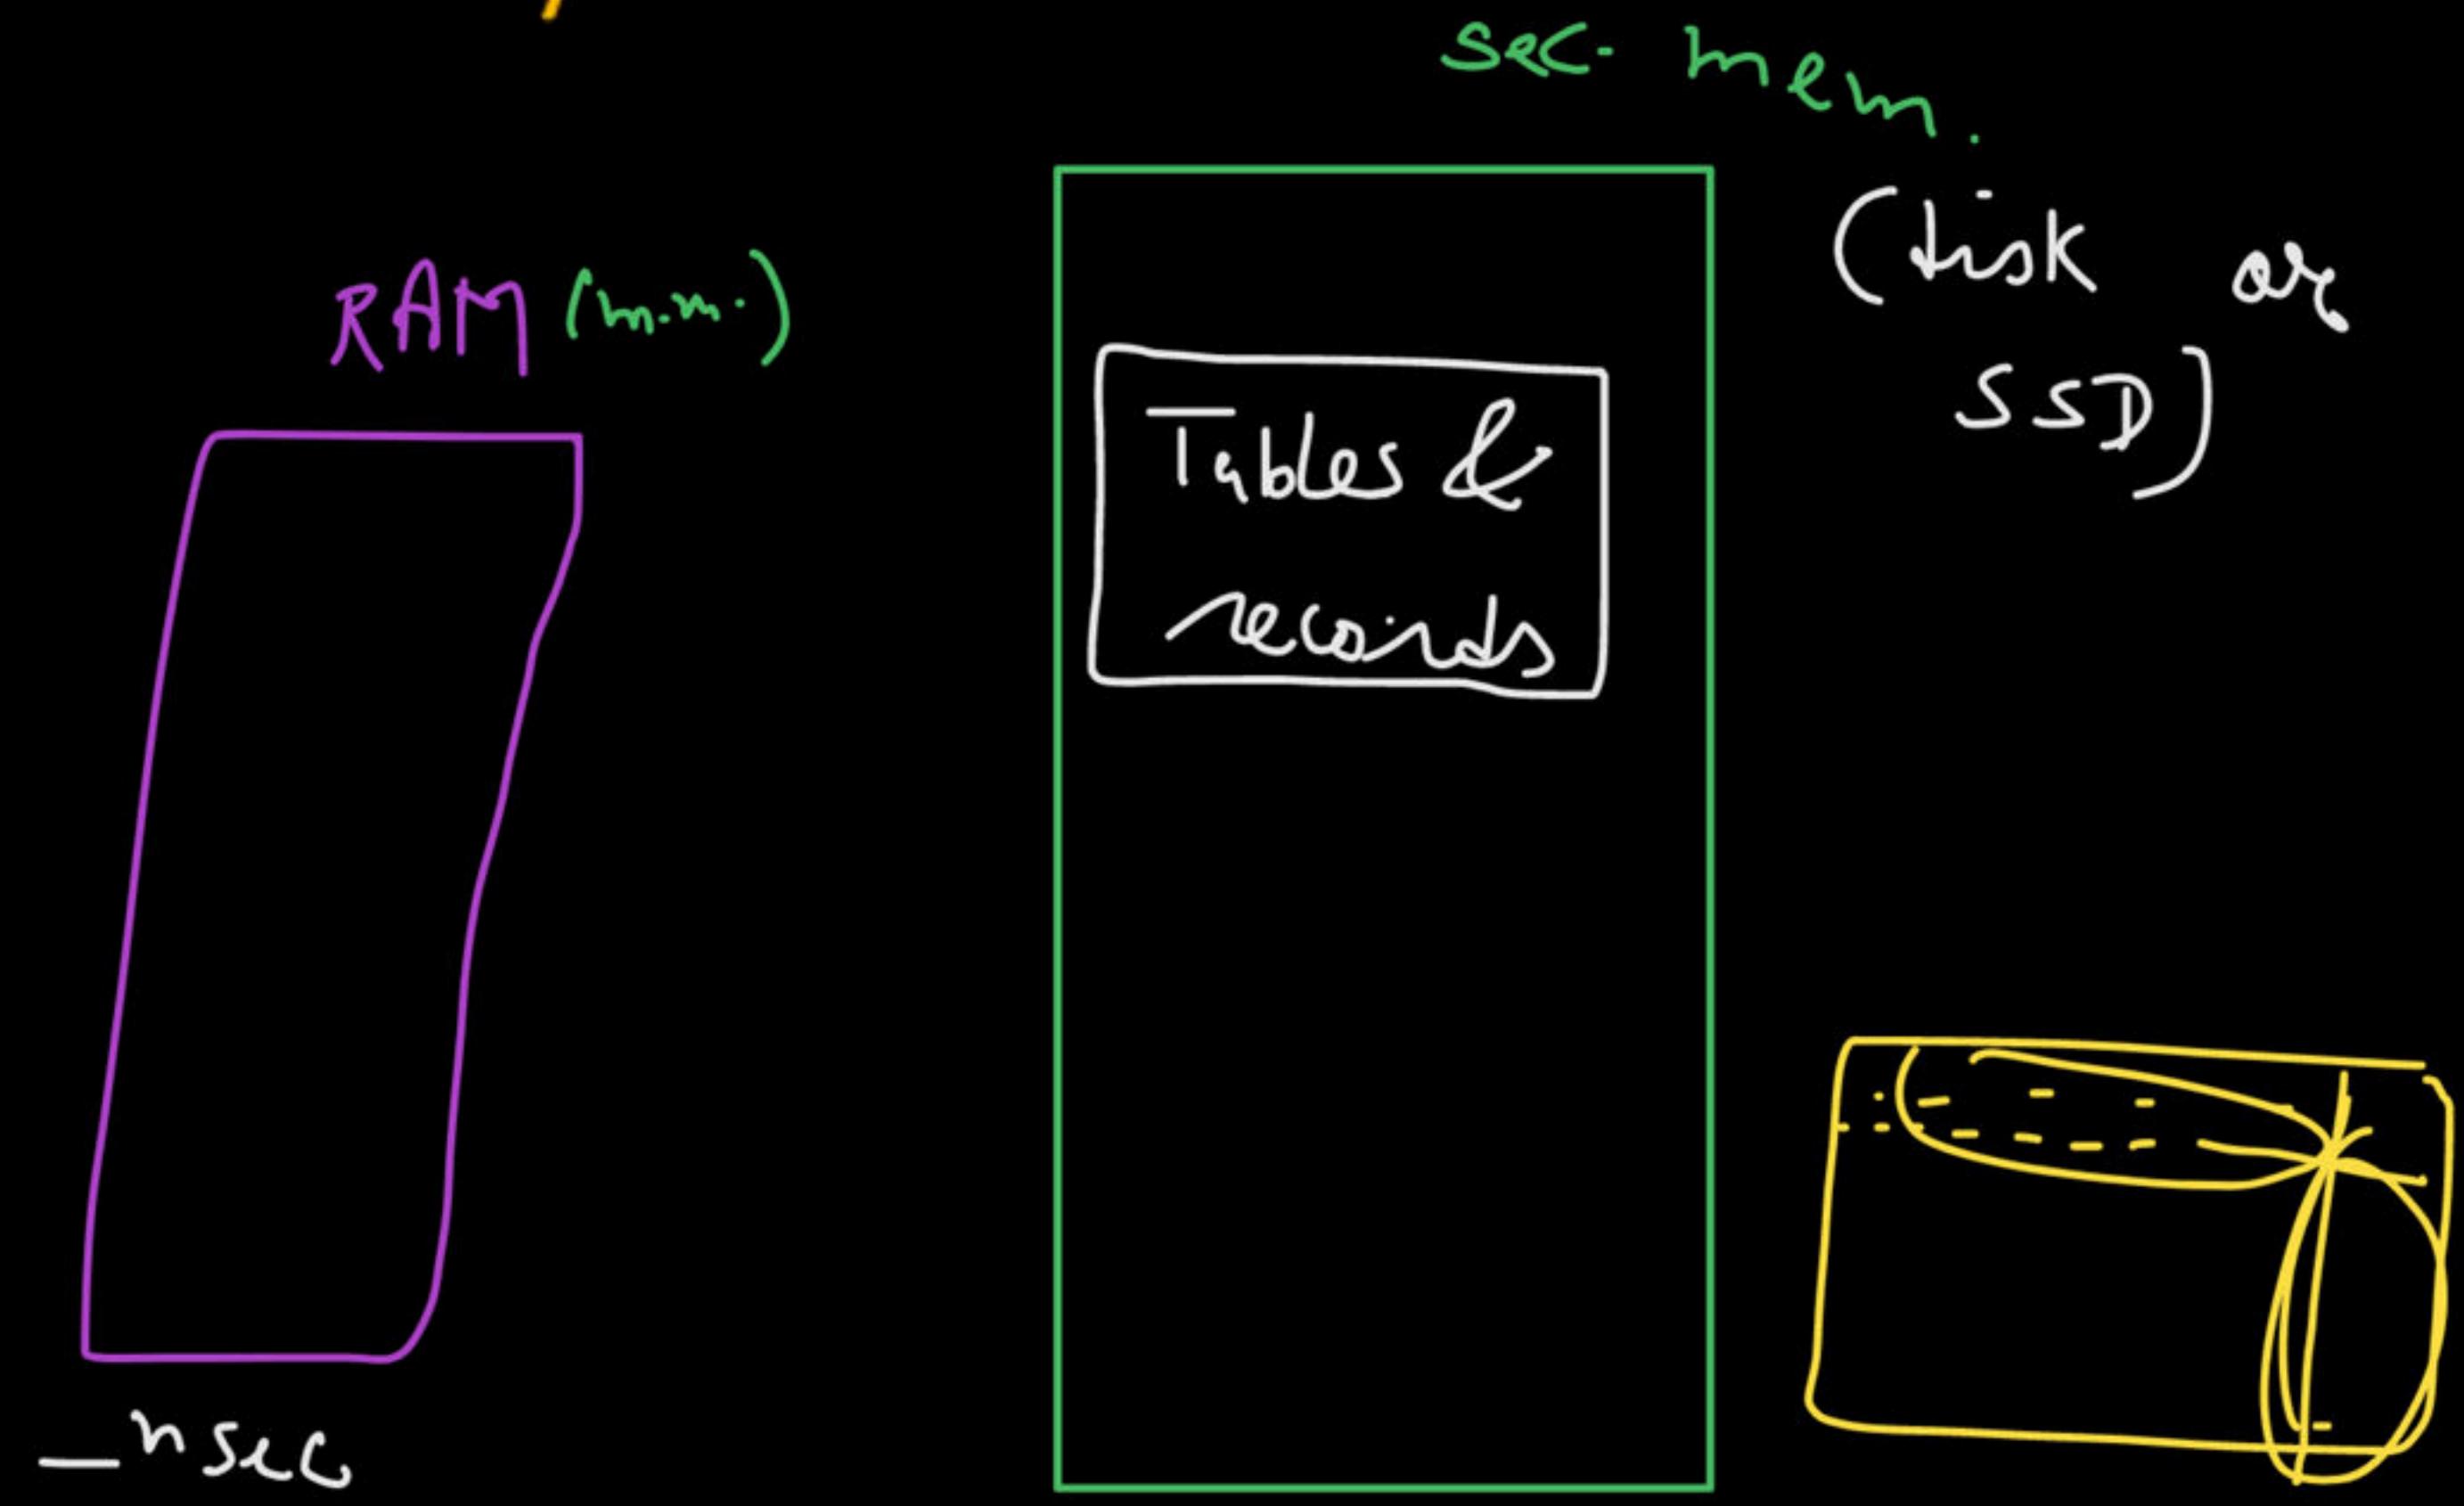
- Multiversion Protocol
- Multigranularity Protocol

Memory Structure

- Main Memory
- Secondary Memory



select * from table



Query optimizer :-

Runs query in such a way that min. time taken to fetch records from disk.

Ex:-

Employees				
Eid	Ename	Gender	Salary	Dateofjoining

select Ename from Employees
where Dateofjoining > 2001 and
Salary > 50000

assume.

n. of records for d.o.b. > 2001 \Rightarrow 5000

n. of records for salary > 50000 \Rightarrow 100

and

both cond' here record \Rightarrow 82

Disk Blocks and Record Storages

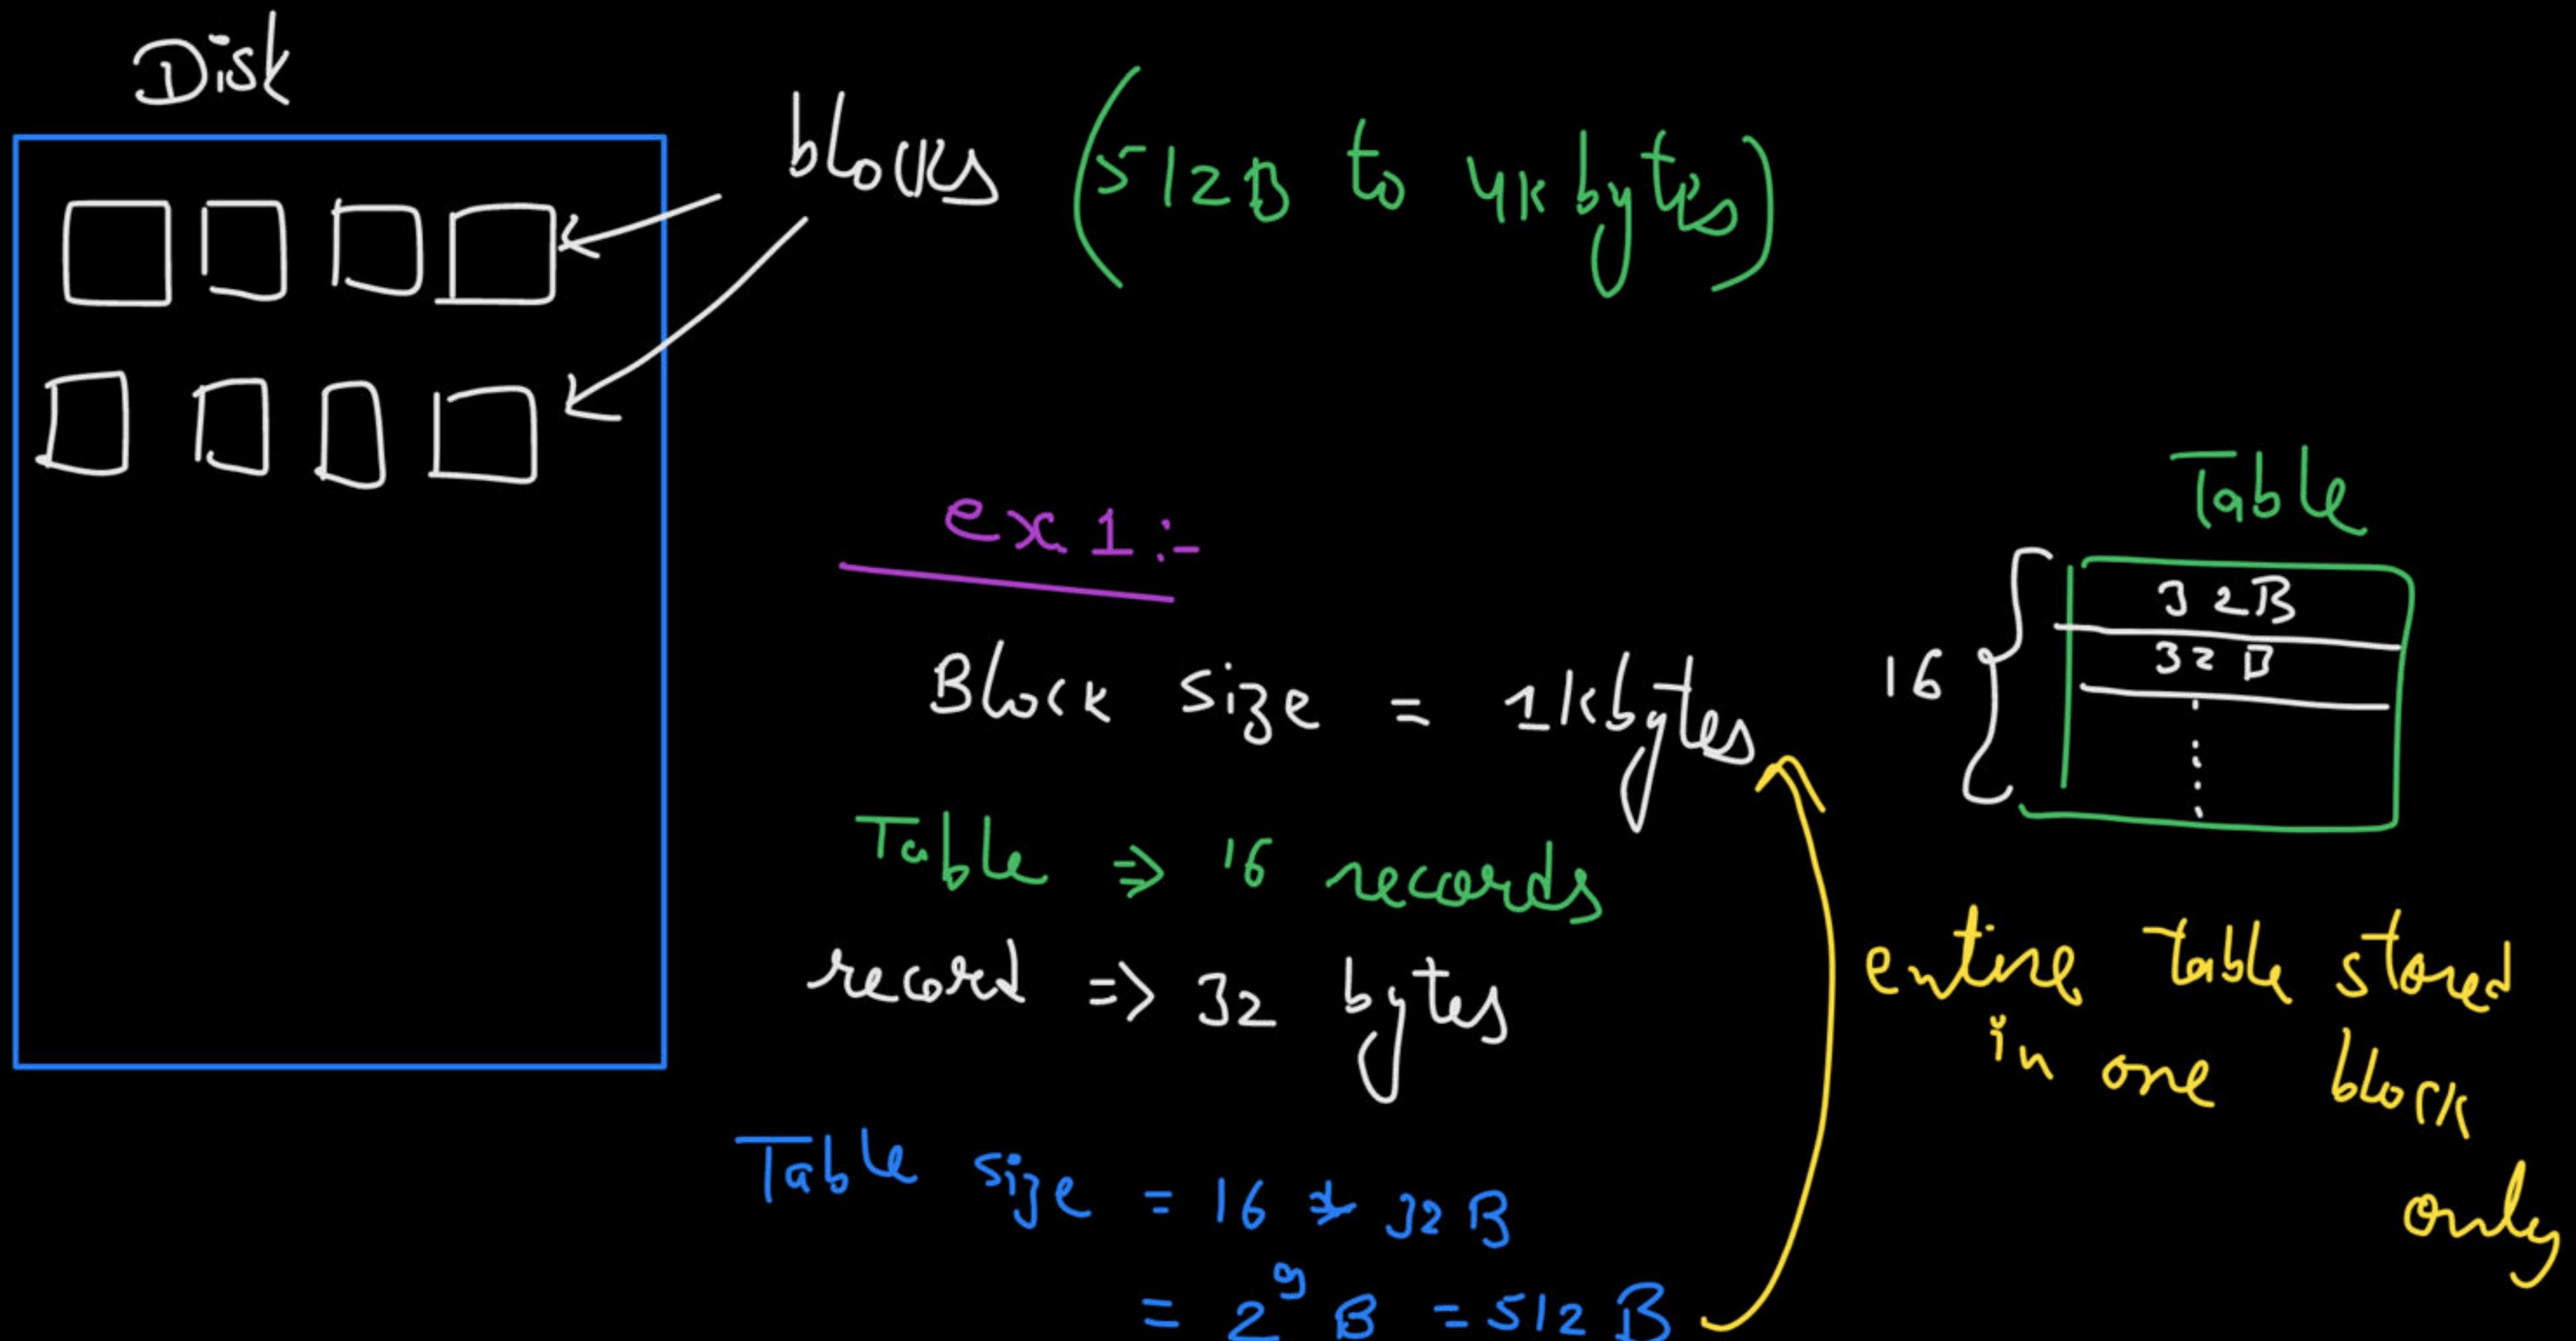


Table \Rightarrow size records

each record \Rightarrow 32 bytes

block size \Rightarrow 1 kbyte

$$\text{Table size} = 512 \times 32B = 2^{14}B = 16KB$$

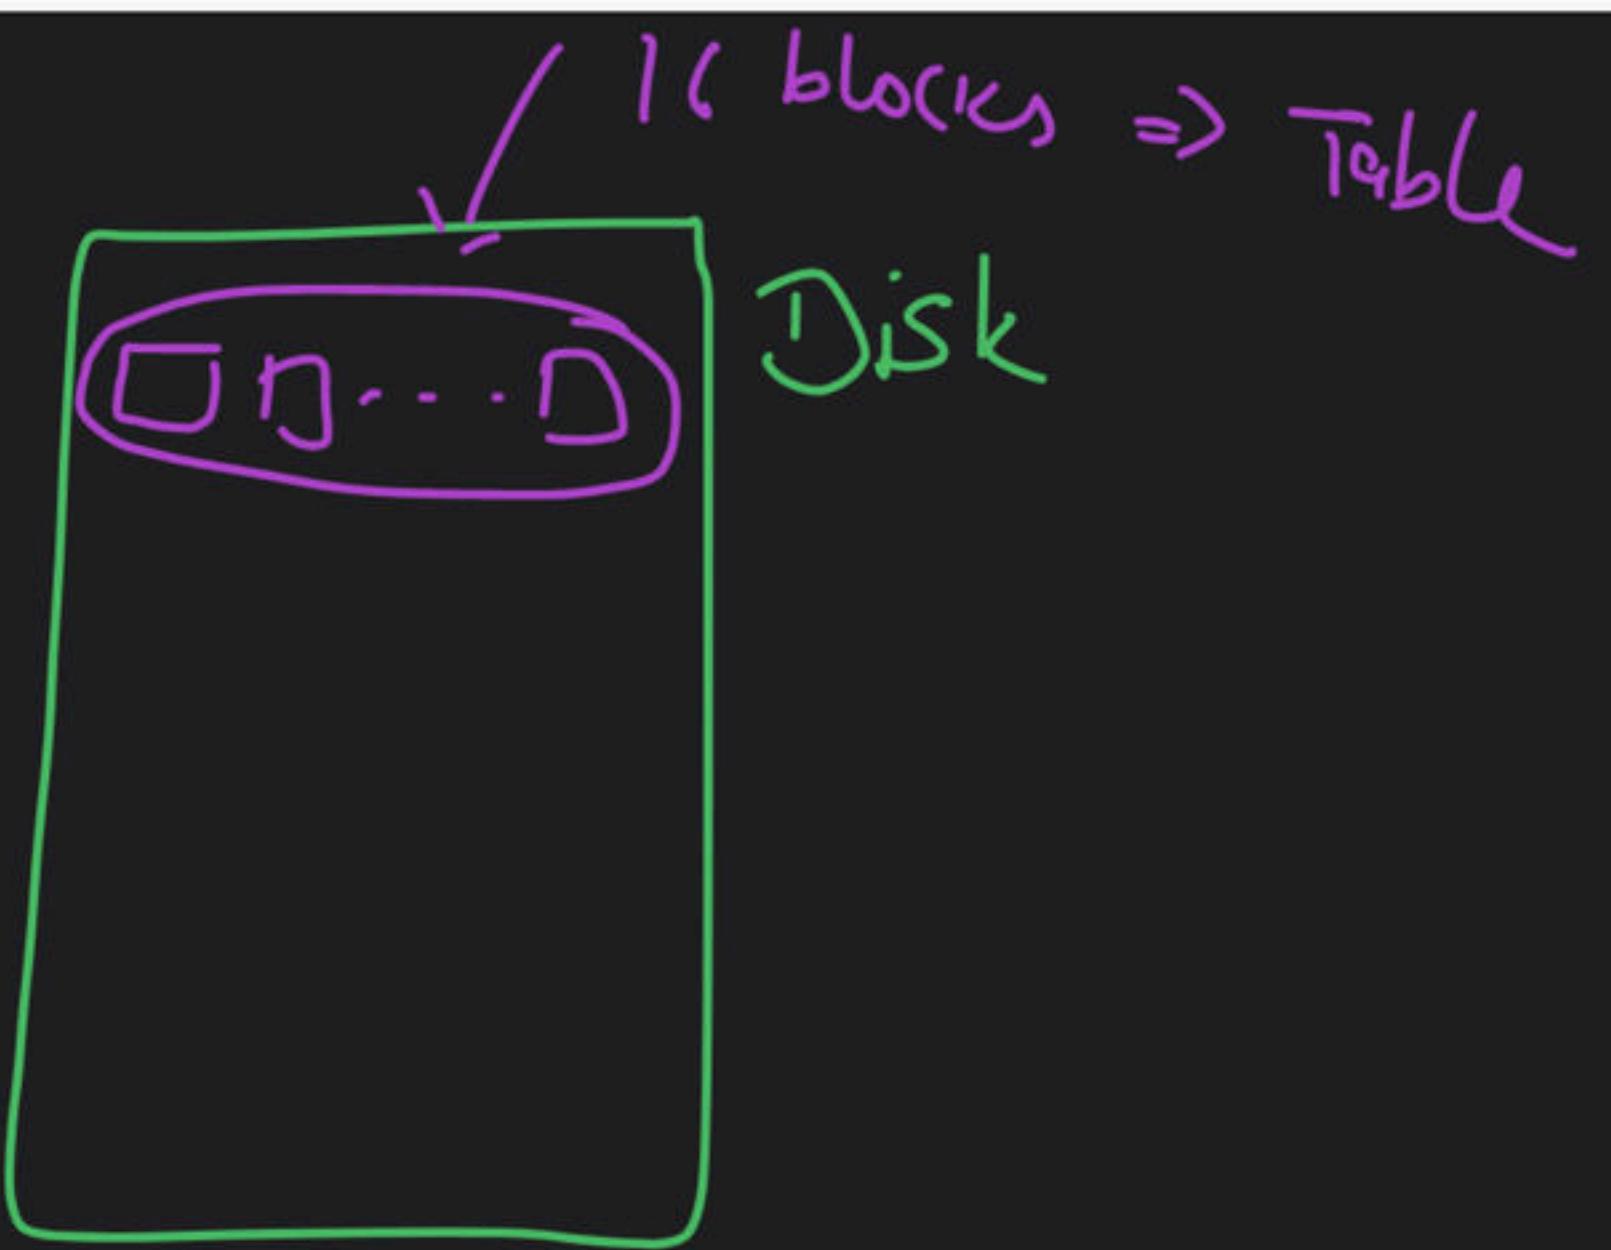
$$\begin{aligned} \text{no. of blocks required to store entire table} &= \frac{16KB}{1KB} = 16 \text{ blocks} \end{aligned}$$

select Ename from employee

where eid = 76

||

Search each record linearly to
get the required record.



Select * From employee where eid < 90



In General all blocks are accessed linearly

but if records are stored in sorted order of eid,
then only till that block where eid becomes 90.

Eid	Ename	Salary
1	E1	57000
2	E2	60000
3	E2	10000
4	E3	45000
5	E4	42000
6	E3	61000
7	E2	55000

Physical storage of db table

- ① Basis on which column the records must be ordered in disk so that max. no. of times queries can provide quick results
- ② Can we reach to a specific record directly.
 - ↳ Through indexing

File Organization

- Sorted File Organization
- Heap File Organization

Fixed Length Records



Fixed Length Records: What if deletion done?

Fixed Length Records: What if deletion done?

Add.	Header Last = 8	Pointer to next free address = 2	
1			
2		5	
3			
4			
5		7	
6			
7		8	
8		NULL	Last
9			

Variable Length Records

Add.	Header Last = 8	Pointer to next free address = 2
1		
2		5
3		
4		
5		7
6		
7		8
8		NULL
9		Last

Why Indexing

Index File

Indexing Techniques

- Clustered Indexing
- Non-Clustered Indexing

Clustered Indexing

- Data order and index order are same

Non-Clustered Indexing

- Data order and index order are not same

Dense Vs Sparse Index

1. Dense: Index record is for each database record
2. Sparse (non-dense): Index record is for a few database records only

Spanned vs Unspanned File Organization

Indexing Techniques

- Primary Indexing
- Clustering Indexing
- Secondary key Indexing
- Secondary non-key Indexing

Primary Indexing

- Indexing done on primary key or any super key
- Data must be ordered on index
- Its always sparse index

Clustering Indexing

- Indexing done on non-key field
- Data must be ordered on index field

Secondary key Indexing

- Indexing done on primary key or any super key
- Data must not be ordered on index field
- It can be dense or sparse index

Secondary Non-key Indexing

- Indexing done on non-key field
- Data must not be ordered on index field

Primary Index

Rno.	Block Pointer
1	B1
5	B2

Rno.	Block Pointer
9	B3
13	B4

1	A
2	B
3	C
4	D
5	A
6	B
7	E
8	W
9	V
10	D
11	I
12	A
13	C
14	S
15	H
16	A

Primary Index

1	A
2	B
3	C
4	D

Select * from Students
where rno=1

5	A
6	B
7	E
8	W

9	V
10	D
11	I
12	A

Primary Index

1	A
2	B
3	C
4	D

Select * from Students
where rno=12

5	A
6	B
7	E
8	W

9	V
10	D
11	I
12	A

Question

Consider a database file of 65536 records, each record of size 64 bytes. Key field is 10 bytes and block pointer size is 22 bytes. Assume that block size is 256 bytes.

1. The number of blocks required to store file?
2. The number of blocks required to store the index file for primary indexing?

Clustering Index

S_Name	Block Pointer
A	B1
B	B2
B	B3
E	B4

1	A	
2	A	B1
3	A	
4	B	
5	B	
6	B	B2
7	B	
8	B	
9	B	
10	B	B3
11	C	
12	D	
13	E	
14	F	B4
15	F	
16	G	

Clustering Index

S_Name	Block Pointer
A	B1
B	B1
C	B3
D	B3
E	B4
F	B4
G	B4

1	A
2	A
3	A
4	B
5	B
6	B
7	B
8	B
9	B
10	B
11	C
12	D
13	E
14	F
15	F
16	G

Question

DB File size 1G records

Record size = 64bytes

Block size = 4096bytes

Index field = 10 bytes

Block pointer size = 22 bytes

Number of distinct values in index file = 16384

Indexing is done on non-key, data is ordered on non-key and indexing is done for each distinct non-key value

Happy Learning.!

